



# Computer Networking: A Top Down Approach



## Chapter 9 Multimedia Networking

7<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson/Addison Wesley  
April 2016



## Chapter 2 Application Layer

8<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson,  
2020

### A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

©anks and enjoy! JFK/KWR

# Video Streaming and CDNs: context


- stream video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube, Amazon Prime: 80% of residential ISP traffic (2020)
- *challenge*: scale - how to reach ~1B users?
- *challenge*: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution*: distributed, application-level infrastructure



# Multimedia: video

- video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- digital image: array of pixels
  - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$




frame  $i+1$

# Multimedia: video

- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- **examples:**
  - MPEG 1 (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, 64Kbps – 12 Mbps)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



frame  $i+1$

# Multimedia networking: 3 example types

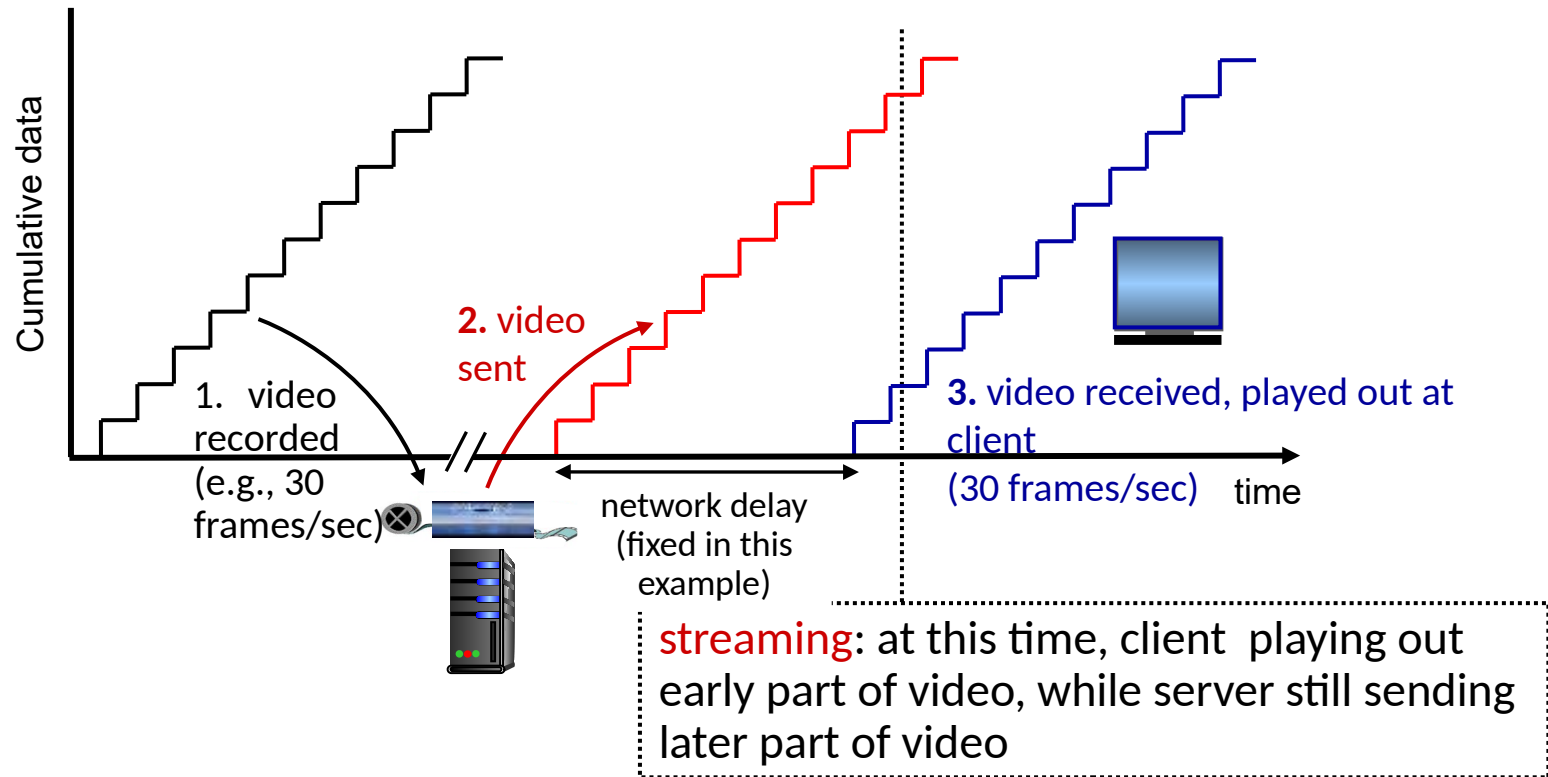
---

- *streaming, stored* audio, video
  - *streaming*: can begin playout before downloading entire file
  - *stored (at server)*: can transmit faster than audio/video will be rendered (implies storing/buffering at client)
  - e.g., YouTube, Netflix, Hulu
- *conversational* voice/video over IP
  - interactive nature of human-to-human conversation limits delay tolerance
  - e.g., Skype
- *streaming live* audio, video
  - e.g., live sporting event (football, hockey, etc.)

# And the future has more




# Streaming stored Video





# Streaming stored Video

simple scenario:



Continuous playout constraint:

- once client playout begins, playback must match original timing
- ... but **network delays are variable** (jitter), so will need **client-side buffer** to match playout requirements

Main challenges:

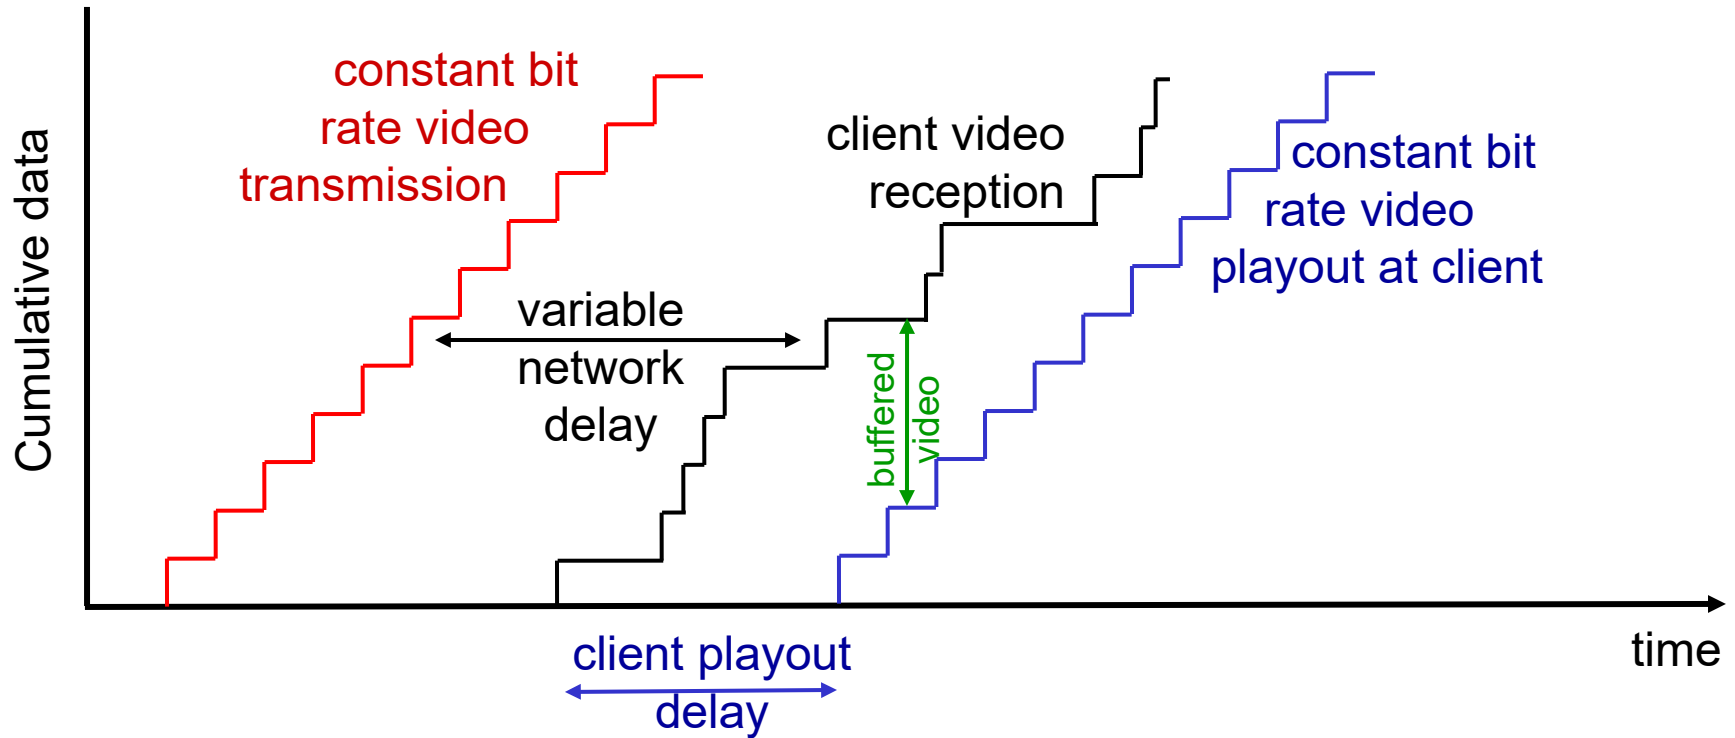
- server-to-client bandwidth will **vary** over time, with changing network congestion levels (in house, access network, network core, video server)
- packet loss, delay due to congestion will delay playout, or result in poor video quality

Other challenges

- client interactivity: pause, fast-forward, rewind, jump through video
- video packets may be lost, retransmitted

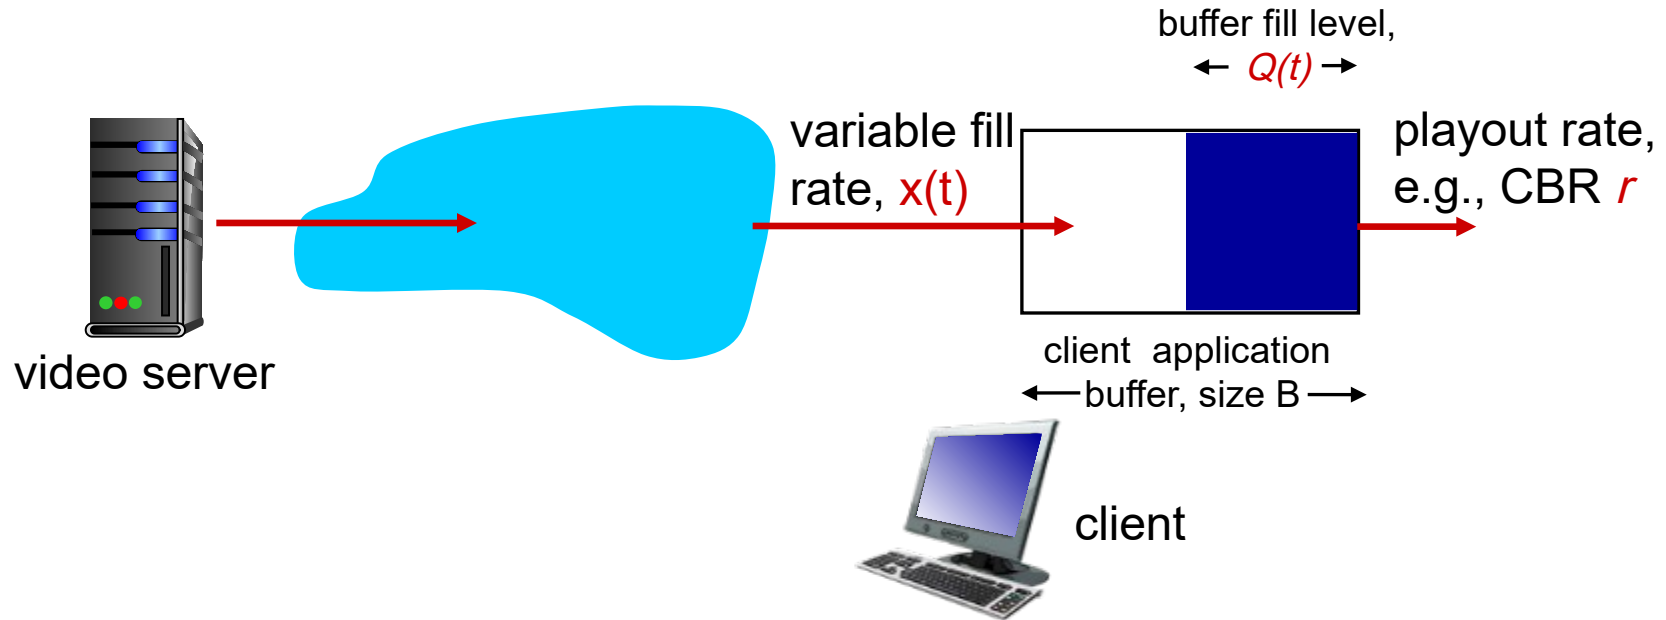


# Streaming stored video: ~~revisited~~

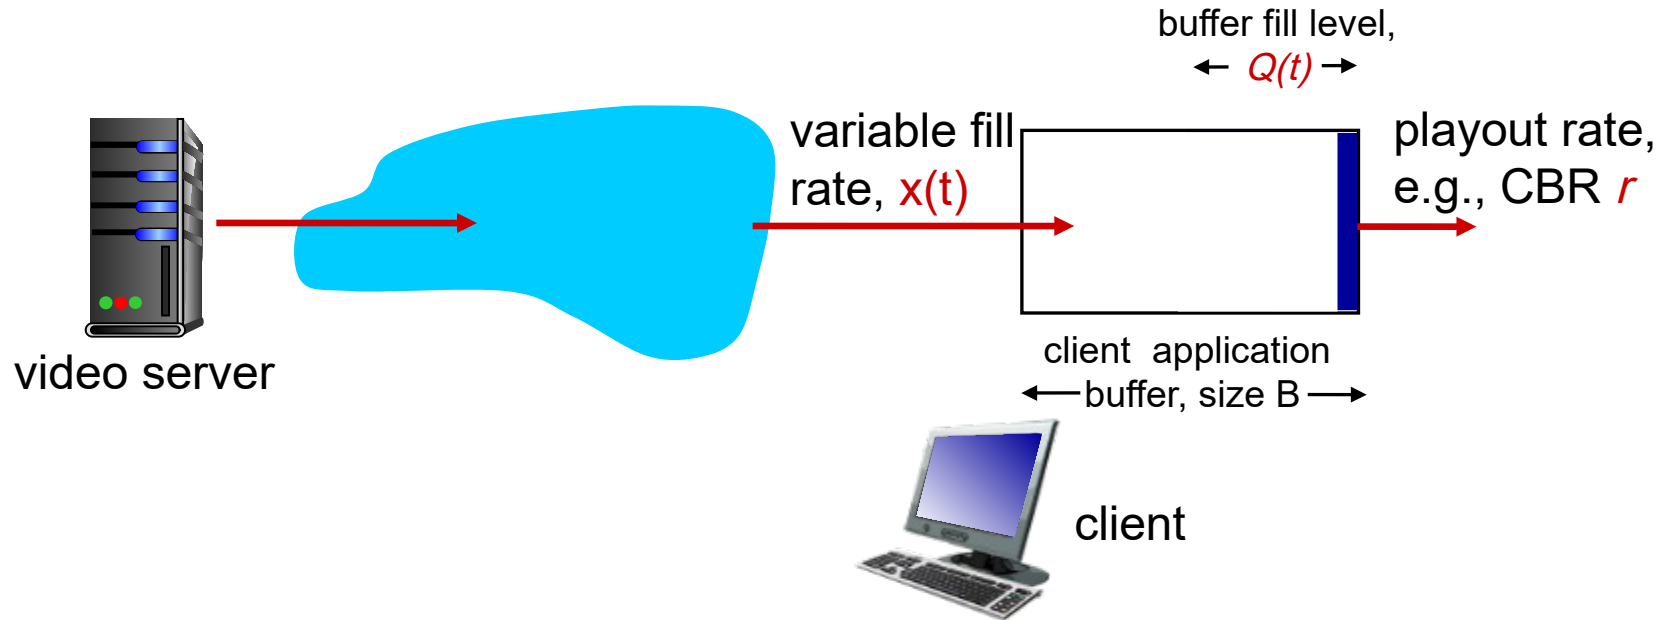


- *client-side buffering and playout delay:* compensate for network-added delay, delay jitter

# Client-side buffering, playout




# Client-side buffering, playout



1. Initial fill of buffer until playout begins at  $t_p$
2. playout begins at  $t_p$ ,
3. buffer fill level varies over time as fill rate  $x(t)$  varies and playout rate  $r$  is constant


# Streaming multimedia: ~~Basic HTTP~~

- multimedia file retrieved via HTTP GET
- send at maximum possible rate under TCP



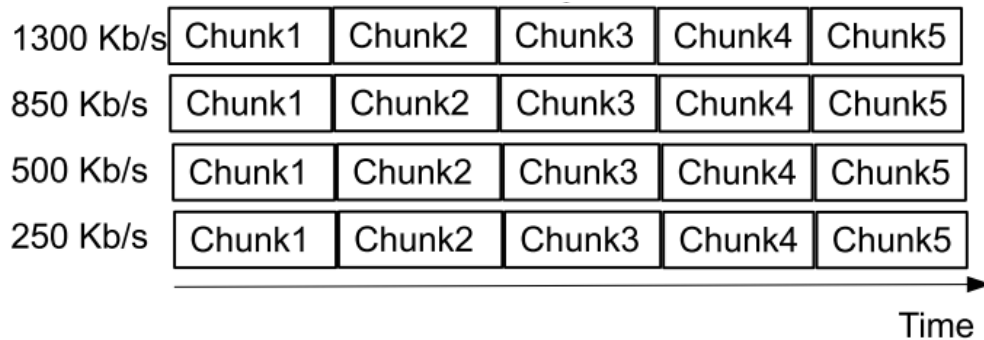
- fill rate fluctuates due to TCP congestion control, retransmissions (in-order delivery)
- larger playout delay: smooth TCP delivery rate

# HTTP-based Adaptive Streaming (HAS) or DASH ...



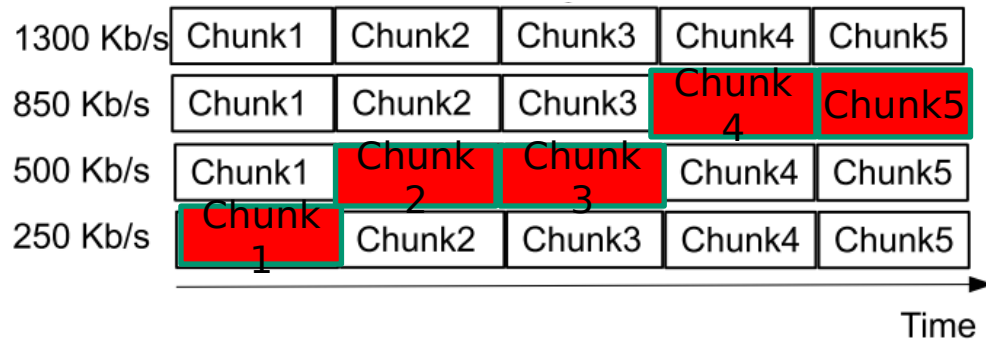
- HTTP-based ~~adaptive~~ streaming
  - Video is split into chunks

# HTTP-based Adaptive Streaming (HAS) or DASH



- HTTP-based **adaptive** streaming
  - Video is split into chunks
  - Each chunk in multiple bitrates (qualities)

# HTTP-based Adaptive Streaming (HAS) or DASH



- HTTP-based **adaptive** streaming
  - Video is split into chunks
  - Each chunk in multiple bitrates (qualities)
  - Clients adapt quality encoding based on buffer/network conditions




# Streaming multimedia: DASH (or HAS)

*D*ynamic, *A*daptive  
*S*teaming over  
*H*TTp

## server:

- divides video file into multiple chunks
- each chunk encoded at multiple different rates
- different rate encodings stored in different files
- files replicated in various CDN nodes
- *manifest file*: provides URLs for different chunks




## client:

- periodically estimates server-to-client bandwidth
- consulting manifest, requests one chunk at a time
  - chooses maximum coding rate sustainable given current bandwidth
  - can choose different coding rates at different points in time (depending on available bandwidth at time), and from different servers

# Streaming multimedia: DASH (or HAS)

- “intelligence” at client: client determines

- *when* to request chunk (so that buffer starvation, or overflow does not occur)
- *what encoding rate* to request (higher quality when more bandwidth available)
- *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)



Streaming video = encoding + DASH + playout buffering

# Content Distribution Networks (CDNs)

*Challenge:* how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- *option 1:* single, large “mega-server”
  - single point of failure
  - point of network congestion
  - long (and possibly congested) path to distant clients

....quite simply: this solution *doesn't scale*

# Content Distribution Networks (CDNs)


*challenge:* how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- *option 2:* store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*)
  - *enter deep:* push CDN servers deep into many access networks
    - close to users
    - Akamai: 240,000 servers deployed in > 120 countries (2015)
  - *bring home:* smaller number (10's) of larger clusters in POPs near access nets
    - used by Limelight



# Content Distribution Networks (CDNs)

- CDN: stores copies of content (e.g. MADMEN) at CDN nodes
- subscriber requests content, service provider returns manifest
  - using manifest, client retrieves content at highest supportable rate
  - may choose different rate or copy if network path congested



# Content Distribution Networks (CDNs)



**OTT challenges:** coping with a congested Internet from the “edge”

- what content to place in which CDN node?
- from which CDN node to retrieve content? At which rate?

# Example: HAS and proxy



*Clients' want*

- ☐ High playback quality
- ☐ Small stall times
- ☐ Few buffer interruptions
- ☐ Few quality switches



# Example: HAS and proxy

Slides from: V. Krishnamoorthi et al.  
"Helping Hand or Hidden Hurdle:  
Proxy-assisted HTTP-based Adaptive  
Streaming Performance",  
Proc. IEEE MASCOIS, 2013



*Clients' want*

- ☐ High playback quality
- ☐ Small stall times
- ☐ Few buffer interruptions
- ☐ Few quality switches

*HAS is increasingly responsible for larger traffic volumes  
... proxies to reduce traffic??*

# Example: HAS and proxy

*Slides from: V. Krishnamoorthi et al.  
"Helping Hand or Hidden Hurdle:  
Proxy-assisted HTTP-based Adaptive  
Streaming Performance",  
Proc. IEEE MASCOIS, 2013*



## *Clients' want*

- ☐ High playback quality
- ☐ Small stall times
- ☐ Few buffer interruptions
- ☐ Few quality switches

## *Network providers' want*

- High QoE of customers/clients

# Example: HAS and proxy

*Slides from: V. Krishnamoorthi et al.  
"Helping Hand or Hidden Hurdle:  
Proxy-assisted HTTP-based Adaptive  
Streaming Performance",  
Proc. IEEE MASCOIS, 2013*



## *Clients' want*


- ☐ High playback quality
- ☐ Small stall times
- ☐ Few buffer interruptions
- ☐ Few quality switches

## *Network providers' want*

- High QoE of customers/clients
- Low bandwidth usage
- High hit rate

# Example: HAS and proxy

*Slides from: V. Krishnamoorthi et al.  
"Helping Hand or Hidden Hurdle:  
Proxy-assisted HTTP-based Adaptive  
Streaming Performance",  
Proc. IEEE MASCOIS, 2013*




## *Clients' want*

- ☐ High playback quality
- ☐ Small stall times
- ☐ Few buffer interruptions
- ☐ Few quality switches

## *Network providers' want*


- High QoE of customers/clients
- Low bandwidth usage
- High hit rate

# Example: HAS and proxy




*Proxy example ...*

# Example: HAS and proxy




# Example: HAS and proxy









# Example: HAS and proxy





# Example: HAS and proxy




# Example: HAS and proxy




# Example: HAS and proxy



*Client 1*






*Proxy before*




*Proxy after*

# Example: HAS and proxy


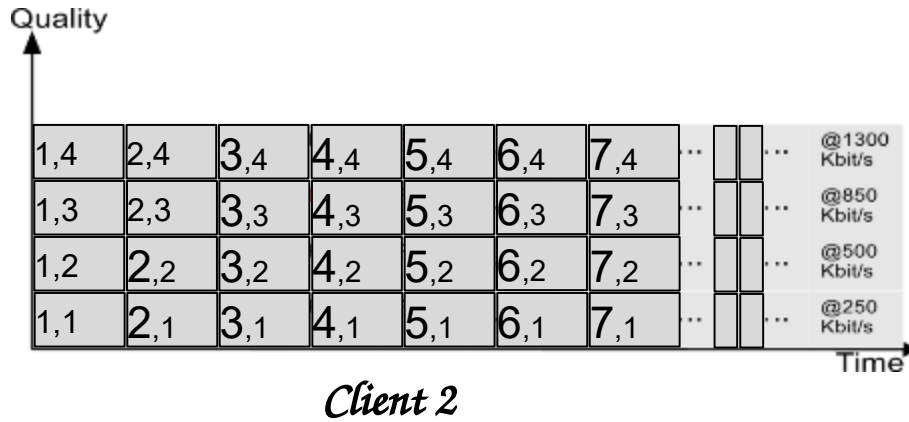



*Proxy before*




*Proxy after*

# Example: HAS and proxy







*Proxy before*





*Proxy after*

# Example: HAS and proxy






# Example: HAS and proxy


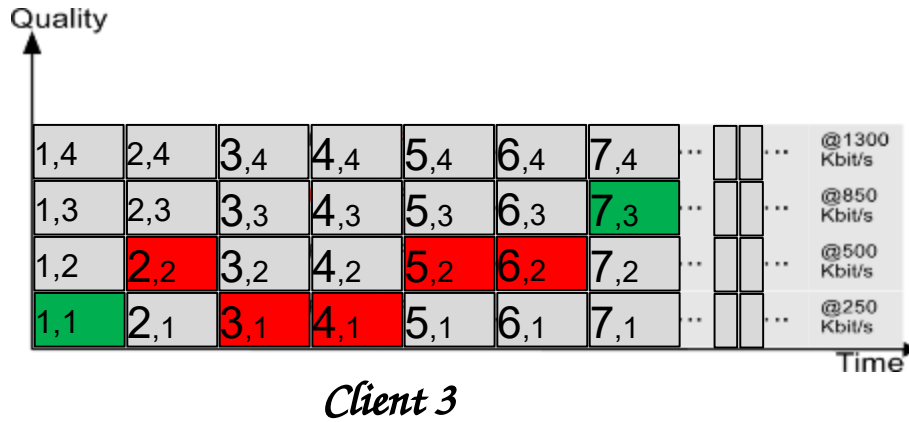



*Proxy before*




*Proxy after*

# Example: HAS and proxy



*Proxy before*



*Proxy after*

# Multimedia networking: ~~outline~~

9.1 multimedia networking  
applications

9.2 streaming *stored* video

9.3 voice-over-IP

9.4 protocols for *real-time*  
conversational applications

9.5 network support for multimedia

# Voice-over-IP (VoIP)

- *VoIP end-end-delay requirement:* needed to maintain “conversational” aspect
  - higher delays noticeable, impair interactivity
  - < 150 msec: good
  - > 400 msec bad
  - includes application-level (packetization, playout), network delays
- *session initialization:* how does callee advertise IP address, port number, encoding algorithms?
- *value-added services:* call forwarding, screening, recording
- *emergency services:* 911


# VoIP characteristics

- speaker's audio: alternating talk spurts, silent periods.
  - 64 kbps during talk spurt
  - pkts generated only during talk spurts
  - 20 msec chunks at 8 Kbytes/sec: 160 bytes of data
- application-layer header added to each chunk
- chunk+header encapsulated into UDP or TCP segment
- application sends segment into socket every 20 msec during talkspurt

# VoIP: packet loss, delay

- *network loss*: IP datagram lost due to network congestion (router buffer overflow)
- *delay loss*: IP datagram arrives too late for playout at receiver
  - delays: processing, queueing in network; end-system (sender, receiver) delays
  - typical maximum tolerable delay: 400 ms
- *loss tolerance*: depending on voice encoding, loss concealment, packet loss rates between 1% and 10% can be tolerated


# Delay jitter



- end-to-end delays of two consecutive packets: difference can be more or less than 20 msec (transmission time difference)

# VoIP: fixed playout delay

- receiver attempts to playout each chunk exactly  $q$  msecs after chunk was generated.
  - chunk has time stamp  $t$ : play out chunk at  $t+q$
  - chunk arrives after  $t+q$ : data arrives too late for playout: data “lost”
- tradeoff in choosing  $q$ :
  - *large  $q$* : less packet loss
  - *small  $q$* : better interactive experience





# Adaptive playout delay (1)

- *goal*: low playout delay, low late loss rate
- *approach*: adaptive playout delay adjustment:
  - estimate network delay, adjust playout delay at beginning of each talk spurt
  - silent periods compressed and elongated
  - chunks still played out every 20 msec during talk spurt
- E.g., adaptively estimate packet delay:  
Exponentially weighted moving average (EWMA)  
[recall TCP RTT estimate]:

*delay estimate  
after  $i$ th packet*

*small constant,  
e.g. 0.1*

*time received - time sent  
(timestamp)  
measured delay of  $i$ th packet*

# VoIP: recovery from packet loss (1)

*Challenge:* recover from packet loss given small tolerable delay between original transmission and playout

- each ACK/NAK takes  $\sim$  one RTT
- alternative: *Forward Error Correction (FEC)*
  - send enough bits to allow recovery without retransmission (recall two-dimensional parity in Ch. 5)

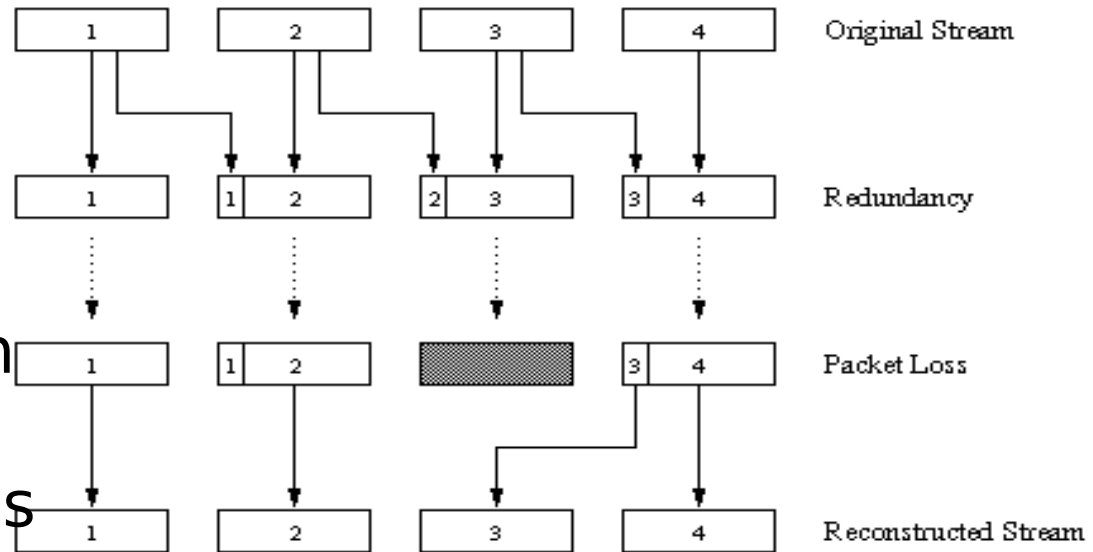
## *simple FEC*

- for every group of  $n$  chunks, create redundant chunk by exclusive OR-ing  $n$  original chunks
- send  $n+1$  chunks, increasing bandwidth by factor  $1/n$
- can reconstruct original  $n$  chunks if at most one lost chunk from  $n+1$  chunks, with playout delay

# VoIP: recovery from packet loss (2)

## another FEC scheme:


- “piggyback lower quality stream”
- send lower resolution audio stream as redundant information
- e.g., nominal stream PCM at 64 kbps and redundant stream GSM at 13 kbps



non-consecutive loss: receiver can conceal loss  
generalization: can also append (n-1)st and (n-2)nd low-bit chunk

# VoIP: recovery from packet loss

## loss (3)




### *interleaving to conceal loss:*

- audio chunks divided into smaller units; e.g., four 5 msec units per 20 msec audio chunk
- packet contains small units from different

- if packet lost, still have *most* of every original chunk
- no redundancy overhead, but increases playout delay

# Voice-over-IP: Skype


- proprietary application-layer protocol (inferred via reverse engineering)
  - encrypted msgs
- P2P components:
  - clients: Skype peers connect directly to each other for VoIP call
  - **super nodes (SN):** Skype peers with special functions
  - **overlay network:** among SNs to locate SCs
  - **login server**



# P2P voice-over-IP: Skype

## Skype client operation:

1. joins Skype network by contacting SN (IP address cached)
2. logs in (username, password) to centralized Skype login server
3. obtains IP address for callee from SN, SN overlay
  - or client buddy list
4. initiate call directly to callee



# Multimedia networking: ~~outline~~

9.1 multimedia networking  
applications

9.2 streaming *stored* video

9.3 voice-over-IP


9.4 protocols for *real-time*  
conversational applications:  
RTP, SIP

9.5 network support for multimedia

# RTP runs on top of UDP

RTP libraries provide transport-layer interface that extends UDP:

- port numbers, IP addresses
- payload type identification
- packet sequence numbering
- time-stamping





# RTP header

|                         |                                 |                   |                                      |                                 |
|-------------------------|---------------------------------|-------------------|--------------------------------------|---------------------------------|
| <i>payload<br/>type</i> | <i>sequence<br/>number type</i> | <i>time stamp</i> | <i>Synchronization<br/>Source ID</i> | <i>Miscellaneous<br/>fields</i> |
|-------------------------|---------------------------------|-------------------|--------------------------------------|---------------------------------|

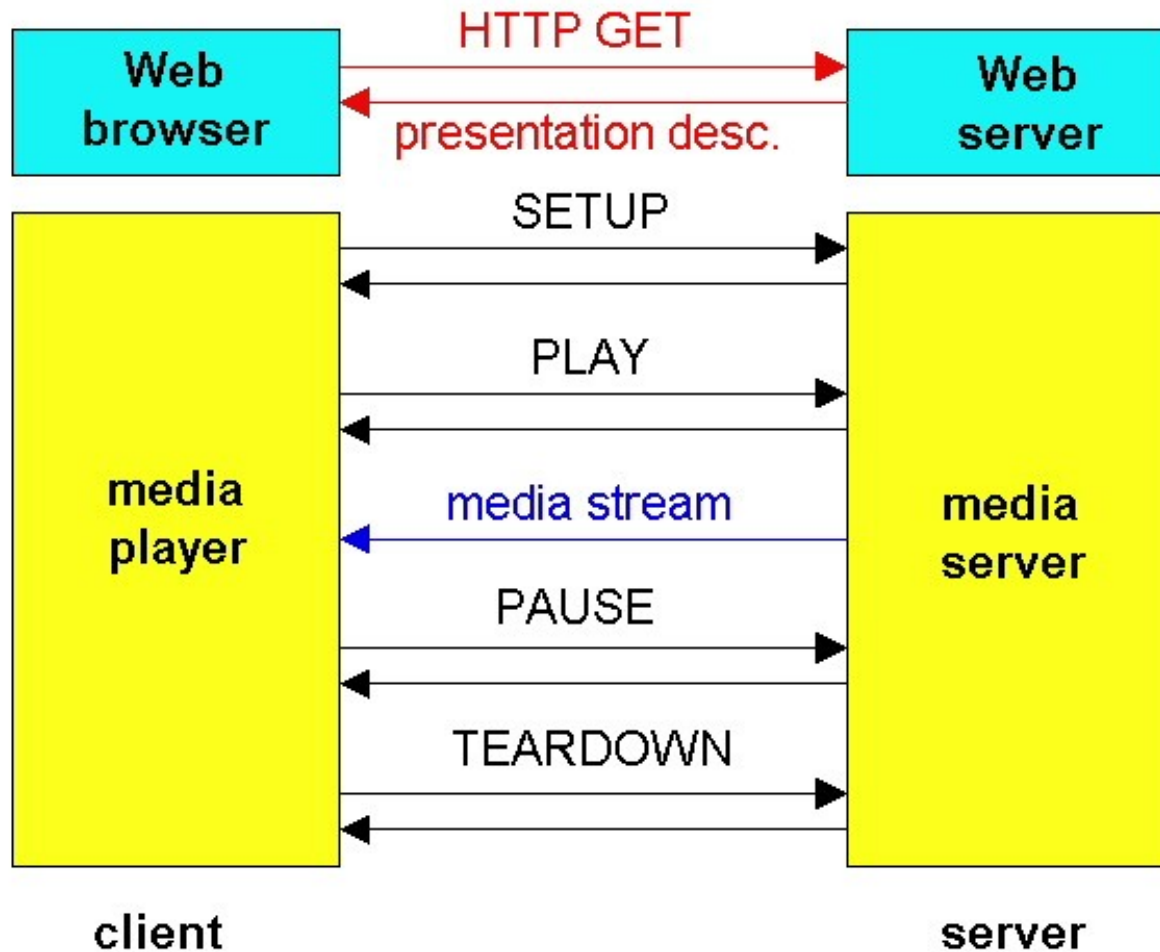
**payload type (7 bits):** indicates type of encoding currently being used. If sender changes encoding during call, sender informs receiver via payload type field

**sequence # (16 bits):** increment by one for each RTP packet sent

- ❖ detect packet loss, restore packet sequence

**timestamp field (32 bits long):** sampling instant of first byte in this RTP data packet (e.g., for audio, timestamp clock increments by one for each sampling period)

# RTSP Operation



Control  
messages  
“out-of-band”  
○ port 554


Media stream  
“in-band”.

# Real-Time Control Protocol (RTCP)

---

- works in conjunction with RTP
- each participant in RTP session periodically sends RTCP control packets to all other participants
- each RTCP packet contains sender and/or receiver reports
  - report statistics useful to application: # packets sent, # packets lost, interarrival jitter
- feedback used to control performance
  - sender may modify its transmissions based on feedback

# RTCP: multiple multicast senders



- E.g., RTCP attempts to limit its traffic to 5% of session bandwidth

## *receiver report packets:*


- fraction of packets lost, last sequence number, average interarrival jitter

## *sender report packets:*

- SSRC of RTP stream, current time, number of packets sent, number of bytes sent

- each RTP session: typically a single multicast address; all RTP /RTCP packets belonging to session use multicast address
- RTP, RTCP packets distinguished from each other via distinct port numbers
- to limit traffic, each participant reduces RTCP traffic as number of conference participants increases

# Fairness of UDP Streams (1/2)




✂  $R1$ - $R2$  is the bottleneck link

✂ Streaming uses UDP at the transport layer; requested media encoded at  $1\text{ Mbps}$

✂ What fraction of the bottleneck is available to FTP?

Credit: MSc thesis work by [Sean Boyden](#) (2006)

# Fairness of RealVideo Streams (2/2)



# Multimedia networking: ~~outline~~

9.1 multimedia networking  
applications

9.2 streaming *stored* video

9.3 voice-over-IP

9.4 protocols for *real-time*  
conversational applications

9.5 network support for multimedia

# Network support for ~~multimedia~~

| Approach                           | Granularity                 | Guarantee                        | Mechanisms  | Complex | Deployed?      |
|------------------------------------|-----------------------------|----------------------------------|---|---------|----------------|
| Making best of best effort service | All traffic treated equally | None or soft                     | No network support (all at application)             | low     | everywhere     |
| Differentiated service             | Traffic “class”             | None or soft                     | Packet market, scheduling, policing.                | med     | some           |
| Per-connection QoS                 | Per-connection flow         | Soft or hard after flow admitted | Packet market, scheduling, policing, call admission | high    | little to none |




# Dimensioning best effort networks


- *approach*: deploy enough link capacity so that congestion doesn't occur, multimedia traffic flows without delay or loss
  - low complexity of network mechanisms (use current “best effort” network)
  - high bandwidth costs
- challenges:
  - *network dimensioning*: how much bandwidth is “enough?”
  - *estimating network traffic demand*: needed to determine how much bandwidth is “enough” (for that much traffic)

# Providing multiple classes of ~~service~~

- thus far: making the best of best effort service
  - one-size fits all service model
- alternative: multiple classes of service
  - partition traffic into classes
  - network treats different classes of traffic differently (analogy: VIP service versus regular service)
- granularity: differential service among multiple classes, **not among individual connections**
- history: ToS bits




# Multiple classes of service: scenario



# Scenario 1: mixed HTTP and VoIP

- example: 1Mbps VoIP, HTTP share 1.5 Mbps link.
  - HTTP bursts can congest router, cause audio loss
  - want to give priority to audio over HTTP




## Principle 1

packet marking needed for router to distinguish between different classes; and new router policy to treat packets accordingly

# Principles for QoS guarantees (more)

- what if applications misbehave (VoIP sends higher than declared rate)
  - policing: force source adherence to bandwidth allocations
- *marking, policing* at network edge




## Principle 2

provide protection (isolation) for one class from others

# Principles for QOS guarantees (more)

- allocating *fixed* (non-sharable) bandwidth to flow: *inefficient* use of bandwidth if flows doesn't use its allocation



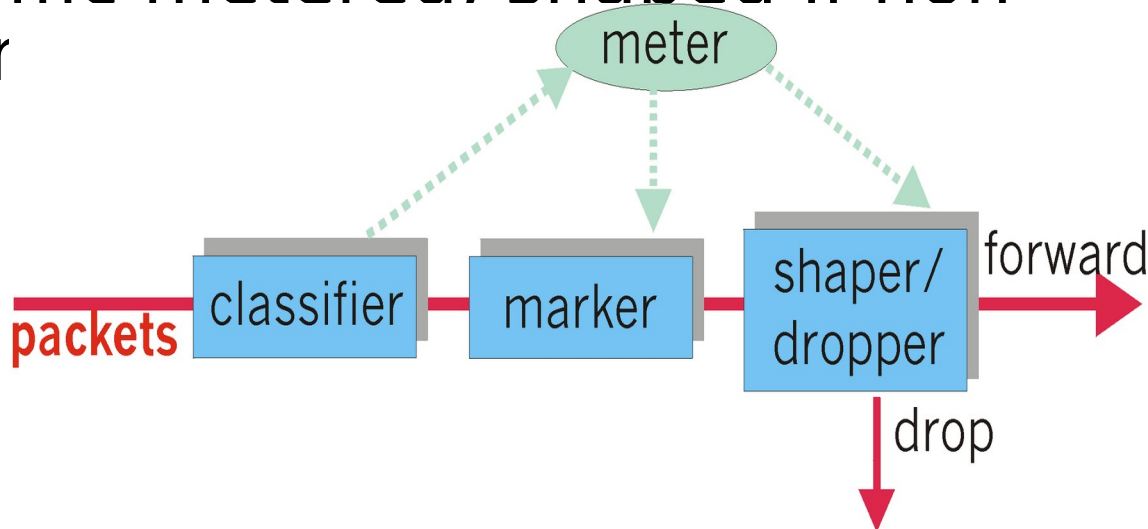
## Principle 3

while providing isolation, it is desirable to use resources as efficiently as possible

# Classification, conditioning


may be desirable to limit traffic injection rate of some class:

- user declares traffic profile (e.g., rate, burst size)
- traffic metered. shaped if non-cor



# Per-connection QoS guarantees

- *basic fact of life*: can not support traffic demands beyond link capacity

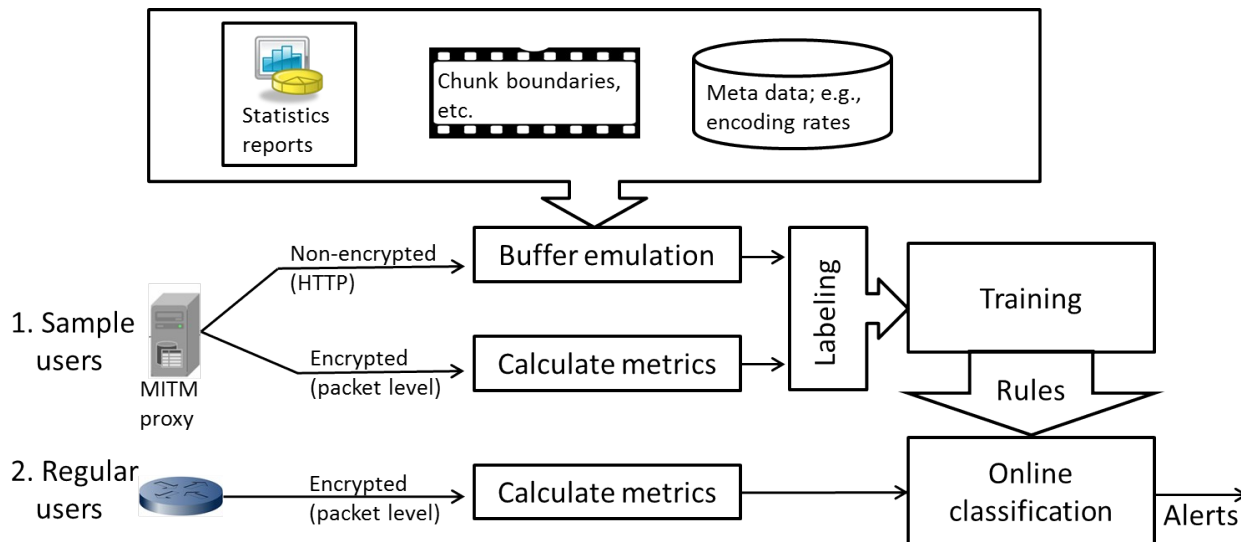



**Principle 4** —  
call admission: flow declares its needs, network may block call (e.g., busy signal) if it cannot meet needs





# ... BUFFEST (per connection discussion)...



ACM MMSys 2017  
(Krishnamoorthi et al.)

# Multimedia networking: outline

---

9.1 multimedia networking  
applications

9.2 streaming *stored* video

9.3 voice-over-IP

9.4 protocols for *real-time*  
conversational applications

9.5 network support for multimedia